

# Halide and Clockwork Compiler

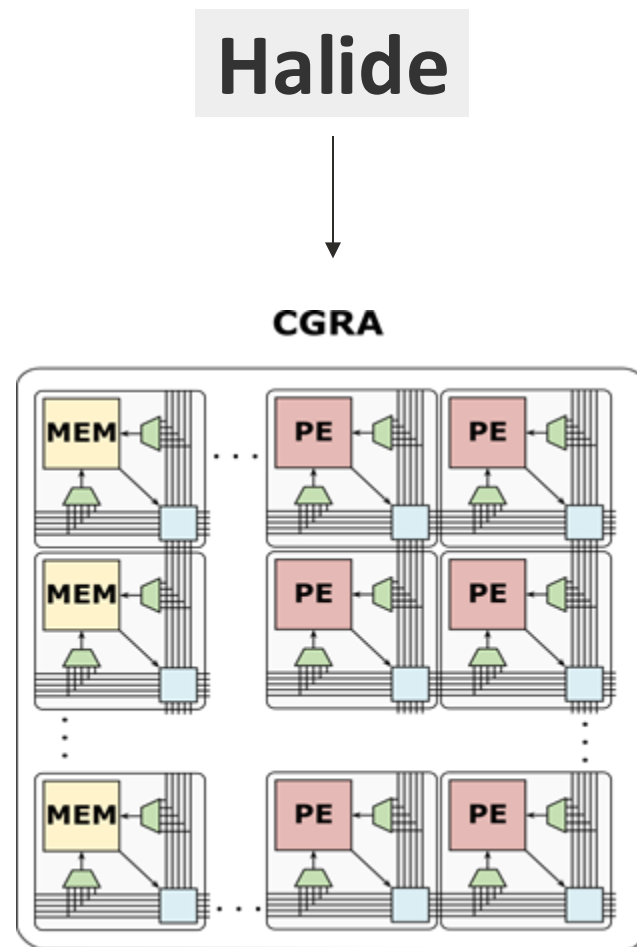
Jeff Setter

# Overview of our System

Applications: in Halide DSL (stencils + DNNs)

Hardware: CGRAs

- Coarse-Grained Reconfigurable Array




# Talk Overview:

## 1) Halide Scheduling

```
app += algorithm;  
app.scheduling();
```

## 2) Clockwork to map Unified Buffers



# Halide

Motivation: need a concise way to describe specializations for each hardware target

=> Separate algorithm from schedule

Algorithm: description of the computation (*what* output values)

- Mathematical operations

Schedule: optimization decisions (*how* to compute)

- Loop optimizations to run fast

Jonathan Ragan-Kelley, Connelly Barnes, Andrew Adams, Sylvain Paris, Fredo Durand, and Saman Amarasinghe.

*Halide: A language and compiler for optimizing parallelism, locality, and recomputation in image processing pipelines.* In PLDI 2013, page 519-530.


# Example Algorithm: cascade for CPU

```
// Algorithm:
// Define 3x3 window of kernel weights
RDom r(0, 3, 0, 3); // Reduction
kernel(0,0) = 1; kernel(0,1) = 2; ... kernel(2,2) = 1;

// First convolution
conv1(x, y) += kernel(r.x, r.y) * input(x + r.x, y + r.y);
conv1_norm(x, y) = conv1(x,y) / 16;

// Second convolution
conv2(x, y) += kernel(r.x, r.y) * conv1_norm(x + r.x, y + r.y);
conv2_norm(x, y) = conv2(x,y) / 16;

output(x, y) = conv2_norm(x, y);
```



# Halide Scheduling: parallelism and memory

Efficient and fast application execution relies on **parallelism** and good **memory management**.

## Optimizations for CPU

- Loop parallelism: unrolling loops, vectorization, threading
- Memory: tiling, memory granularity (fusion)

# Halide Scheduling: parallelism and memory

Efficient and fast application execution relies on **parallelism** and good **memory management**.

## Optimizations for CPU

- Loop parallelism: unrolling loops, vectorization, threading
- Memory: tiling, memory granularity (fusion)

## Scheduling for **CGRAs**

- Hardware parallelism: PEs simultaneously executing
- Memory: tiling, memory hierarchy, specialized for streaming

Jeff Setter. *Compiling Image Processing and Machine Learning Applications to Reconfigurable Accelerators*. In Stanford Doctoral Dissertation, 2023.

# Halide Scheduling to CGRA

Reused the existing Halide compiler toolchain as much as possible without introducing new primitives.

Hardware scheduling:

- Accelerator: input, output, image size
- Loops: reorder loops, tiling
- Memory: memory temporaries, hierarchy
- Compute: duplication



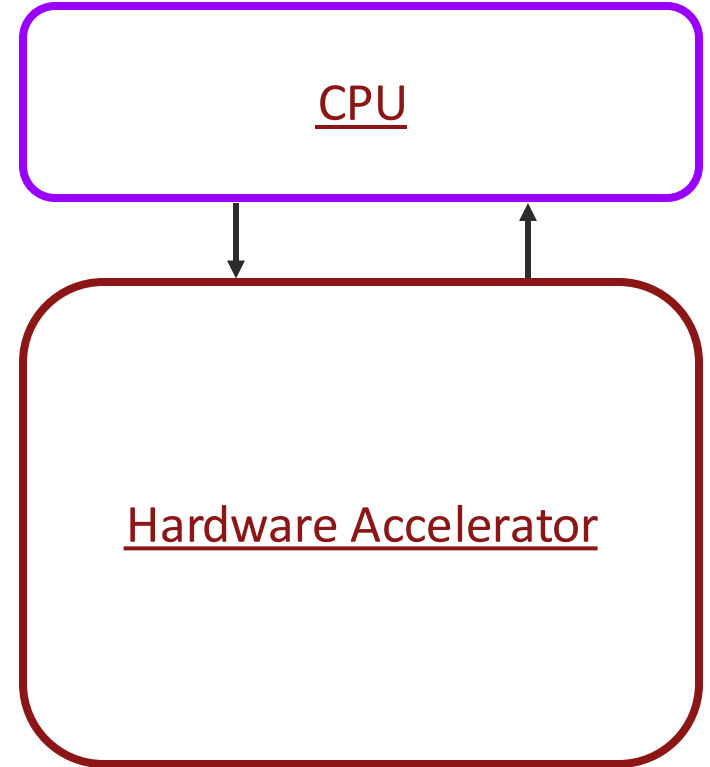
# Define Accelerator's Scope

`hw_accelerate(xi, xo)`

- Defines the accelerator output
- Argument specifies outermost loop

`stream_to_accelerator()`

- Defines an accelerator input



# Define Loops: strip mining, reordering, tiling

```
split(x, xo, xi, 64)
```

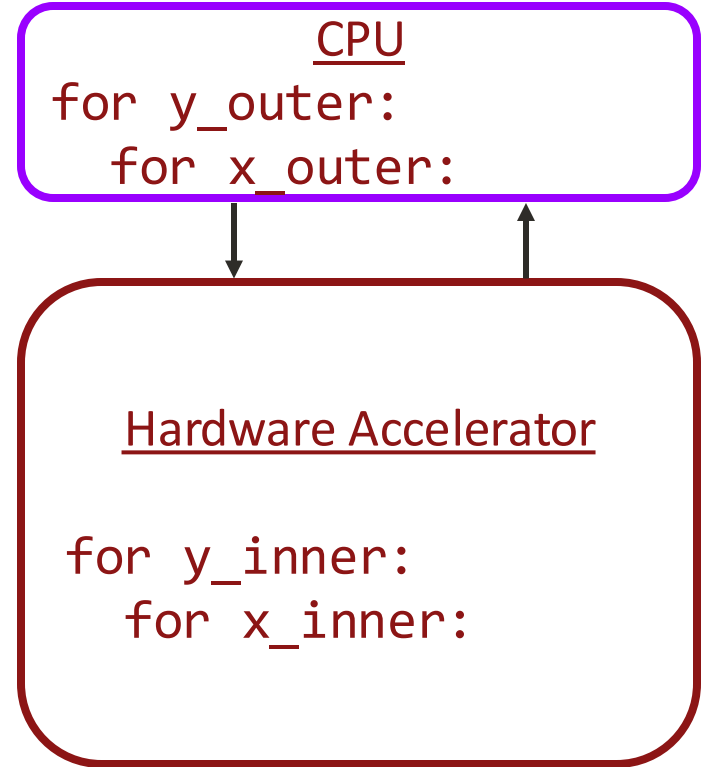
- create nested loops
- specified size for inner loop

```
reorder(xi, yi, xo, yo)
```

- new loop order (from inner to outer)

```
tile(x,y, xo,yo, xi,yi, 64,64)
```

- split two loops and reorder
- common for images



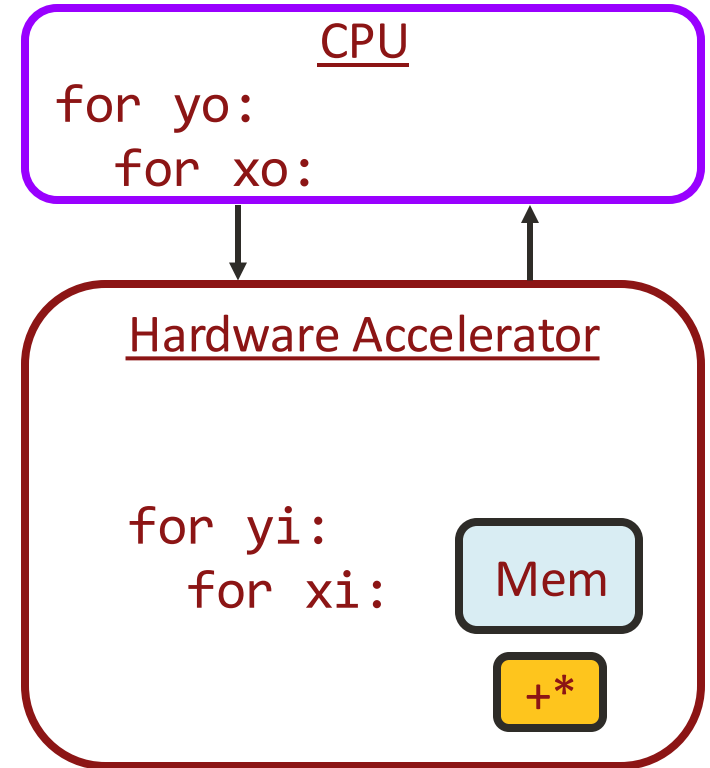
# Define Memories: temporaries

store\_at(consumer, xo)

- create buffers between compute
- at which loop level to create buffer

compute\_at(consumer, xo)

- for CGRA, we compute at tile level and leave loop fusion for Clockwork
- not: choosing line buffer or double buffer



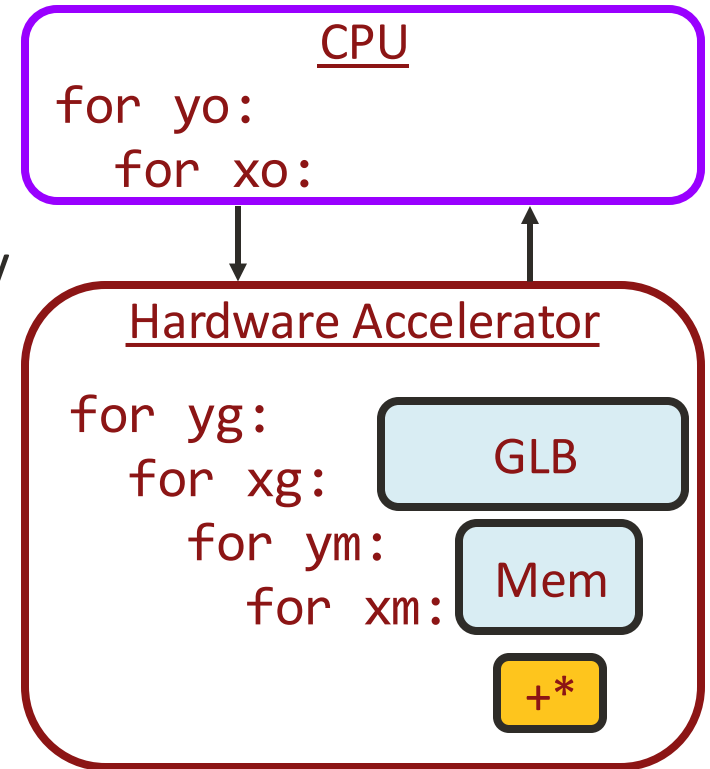
# Define Memories: hierarchy

`in()`

- creates a copy (identity function)
- Needed to move data between memory hierarchy levels with `tile` and `compute_at`

`store_in(MemoryType::GLB)`

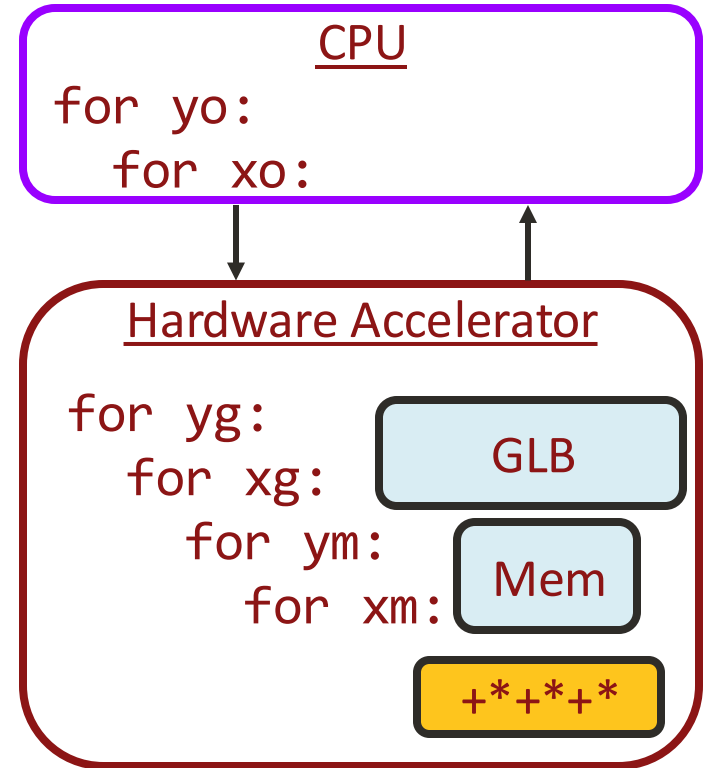
- specify what storage type to use
- Examples: Host, GLB, Memory, Pond



# Define Compute: duplication and sharing

```
unroll(r.y, 3)
```

- increases hardware parallelism
- in Halide IR duplicates loop, because multiple statements = utilizes more existing hardware



# CGRA schedule - cascade

```
1  /* CGRA Schedule */
2  Var xii, xio, yii, yio, xi, xo, yi, yo;
3  output.bound(x, 0, outImgSizeX)
4      .bound(y, 0, outImgSizeY);
5
6  // Accelerate 360x360 tiles in GLB
7  hw_output.in().compute_root()
8      .tile(x,y, xo,yo, xi,yi, 360,360)
9      .hw_accelerate(xi, xo);
10
11 // Send 60x60 tiles to CGRA fabric
12 hw_output
13     .tile(x,y, xio,yio, xii,yii, 60,60)
14     .compute_at(hw_output.in(), xo)
15     .store_in(MemoryType::GLB);
16
17 // conv2 kernel with unrolled reduction
18 conv2_norm.compute_at(hw_output, xio);
19 conv2.update()
20     .unroll(r2.x).unroll(r2.y);
21
22 // conv1 kernel with unrolled reduction
23 conv1_norm.compute_at(hw_output, xio);
24 conv1.update()
25     .unroll(r.x).unroll(r.y);
26
27 // MemoryTile and GLB for input stream
28 hw_input.in().in()
29     .compute_at(hw_output, xio);
30 hw_input.in()
31     .compute_at(hw_output.in(), xo)
32     .store_in(MemoryType::GLB);
33 hw_input.accelerator_input();
34
35 kernel.compute_at(hw_output, xio);
```

CPU - output

Hardware Accelerator

GLB

conv2

Mem

conv1

Mem

GLB

CPU - input

# Halide Compiler and Codegen

- Compiler: performs loop transformations
- Codegen to Clockwork: generates files of **compute kernels** and buffers defined as **sequences of memory accesses**


```
1 // Cascade has a unified buffer named "conv2_stencil" storing 16-bit values
2 prg.buffer_port_widths["conv2_stencil"] = 16;
3
4 // These loops correspond to the iteration domain
5 auto y = prg.add_loop("y", 0, 99);
6 auto x = y->add_loop("x", 0, 600);
7
8 // The function indicates which compute kernel to use
9 auto hcompute_conv2_stencil_1 = x->add_op("op_hcompute_conv2_stencil_1");
10 hcompute_conv2_stencil_1->add_function("hcompute_conv2_stencil_1");
11
12 // These loads are access maps for outputs of the "conv1_stencil" buffer
13 hcompute_conv2_stencil_1->add_load("conv1_stencil", "y", "x");
14 hcompute_conv2_stencil_1->add_load("conv1_stencil", "y", "(x + 1)");
15 hcompute_conv2_stencil_1->add_load("conv1_stencil", "y", "(x + 2)");
16 hcompute_conv2_stencil_1->add_load("conv1_stencil", "(y + 1)", "x");
17 hcompute_conv2_stencil_1->add_load("conv1_stencil", "(y + 1)", "(x + 1)");
18 hcompute_conv2_stencil_1->add_load("conv1_stencil", "(y + 1)", "(x + 2)");
19 hcompute_conv2_stencil_1->add_load("conv1_stencil", "(y + 2)", "x");
20 hcompute_conv2_stencil_1->add_load("conv1_stencil", "(y + 2)", "(x + 2)");
21 hcompute_conv2_stencil_1->add_load("conv1_stencil", "(y + 2)", "(x + 1)");
22 hcompute_conv2_stencil_1->add_load("conv2_stencil", "y", "x");
23
24 // This store is an access map for the input of the "conv2_stencil" buffer
25 hcompute_conv2_stencil_1->add_store("conv2_stencil", "y", "x");
```

# Talk Overview:

## 1) Halide Scheduling

```
app += algorithm;  
app.scheduling();
```

## 2) Clockwork to map Unified Buffers

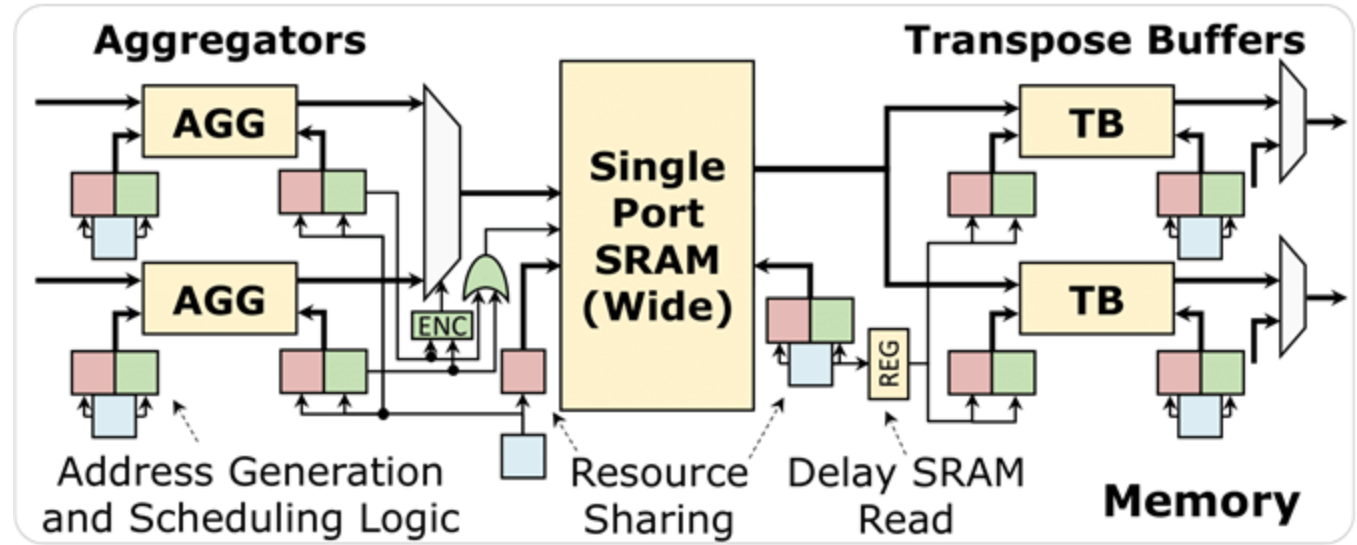




# Unified Buffer Motivation

## Motivation:

need abstraction  
for memories in  
streaming  
accelerators

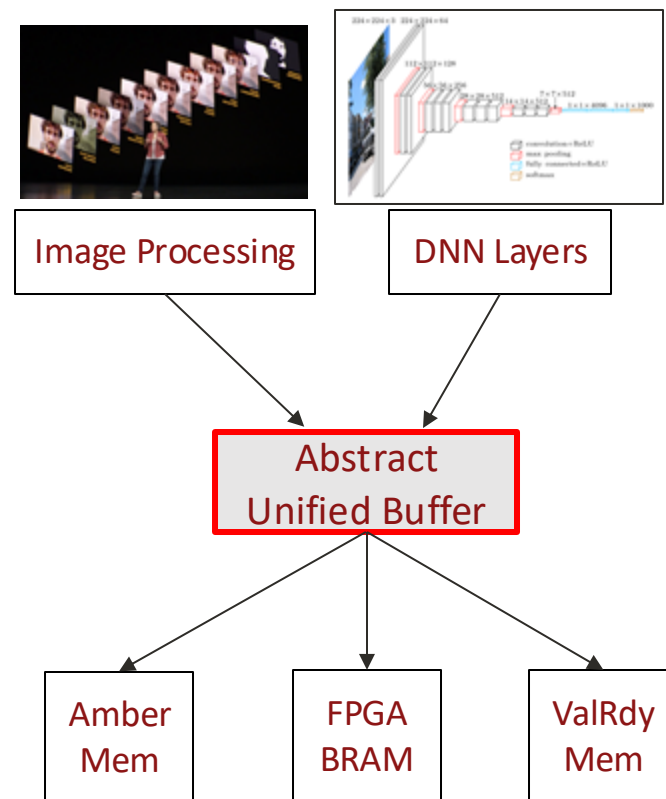


Push memories are **efficient**, but also **complex**

# Unified Buffer Abstraction

Create an abstraction to **interface** between the Halide **application** and the **hardware** motifs

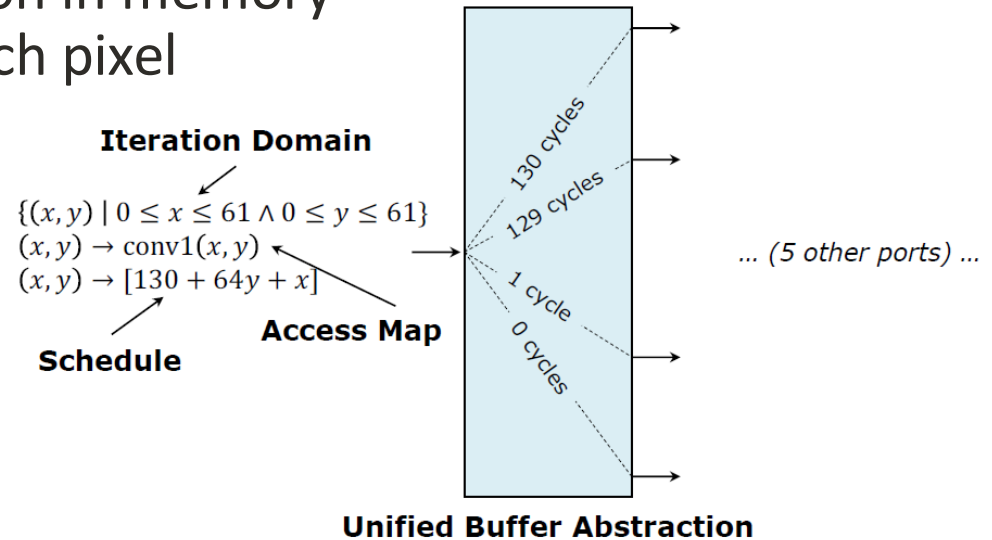
- Domain: fully describes streaming applications (stencils, DNNs)
- Hardware Mapping:
  - Sufficiently general to map to dataflow accelerators classes (FPGA, CGRA; statically scheduled, ready-valid)
  - with efficiency of common hardware stream memories (linebuffers, double buffers)



# Unified Buffer Properties

Properties for each memory port:

- dependencies: relation to other memories
- iteration domain: valid indices
- access map: relative location in memory
- op schedule: cycle time each pixel executes



# Task for Memory Mapper: Clockwork

## Optimize

- Improve schedule: reduce latency as much as possible using loop fusion and loop pipelining


## Map to memories

- Occurs after schedule optimization
- Account for hardware constraints

## Implemented by Qiaoyi



Qiaoyi Liu. *Compiling Applications to Reconfigurable Push Memory Accelerators*. In Stanford Doctoral Dissertation, 2023.

# Unified Buffer: Clockwork scheduling




## Clockwork scheduling:

- Polyhedral analysis in Clockwork, created by Dillon + Qiaoyi
- Purpose: statically determine cycle times when all stores/loads occur
- Calculation: determined by dependencies, then minimize latency
- Buffer size: Number of cycles between store and last load for each pixel

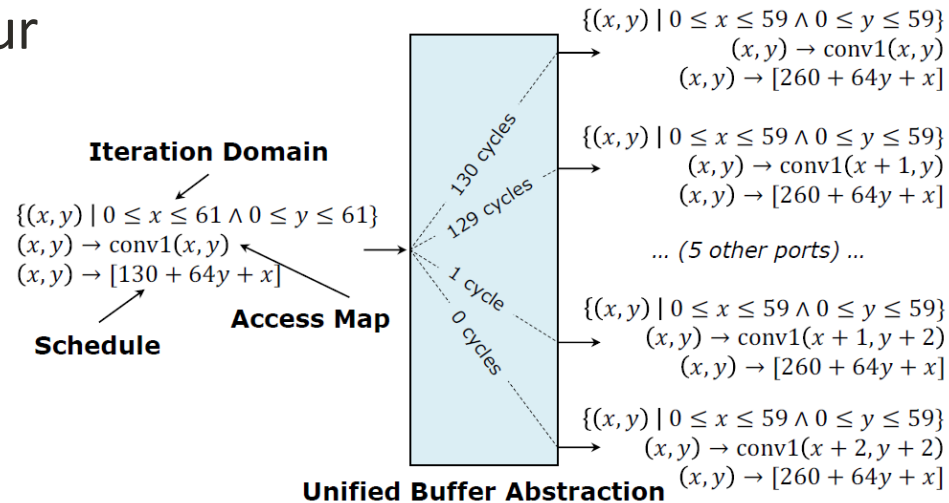


# Unified Buffer: Clockwork scheduling



## Clockwork scheduling:

- Polyhedral analysis in Clockwork, created by Dillon + Qiaoyi
- Purpose: statically determine cycle times when all stores/loads occur
- Calculation: determined by dependencies, then minimize latency
- Buffer size: Number of cycles between store and last load for each pixel



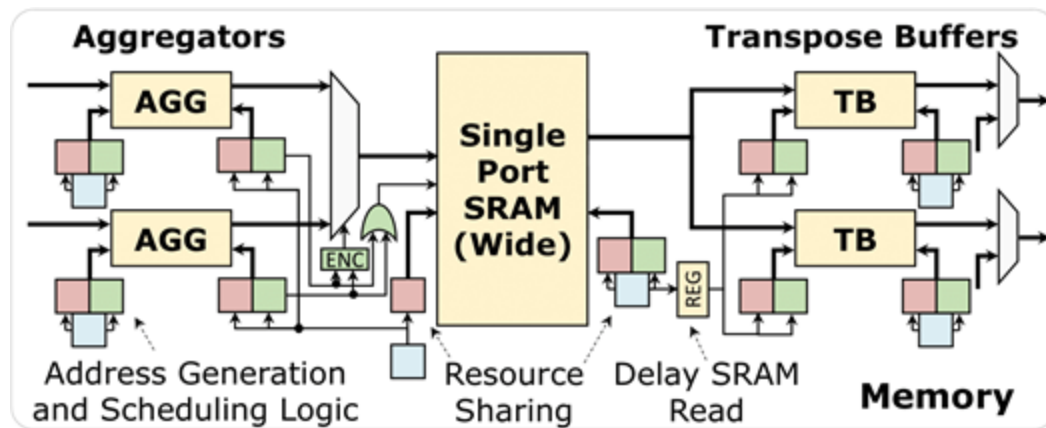
# Unified Buffer: mapping to MEM tiles

## Mapping to hardware

- Occurs after unified buffer abstraction and scheduling
- **One unified buffer != one memory tile**

## Amber CGRA Memory Tile:


- 2048 B capacity
- 2 inputs, 2 outputs
- 6-level affine loopnest
- 4-wide SRAM reads/writes



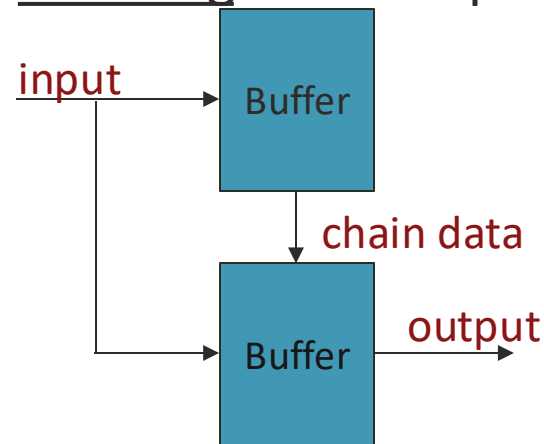
# Clockwork: mapping steps

Key steps:

- Wide fetch: align with 4-wide single-port SRAM
- Banking: more IOs
- Shift register optimization: identify locality opportunities



- Chaining: more capacity






# Summary

```
app += algorithm;  
app.scheduling();
```

Used **Halide** to specify an application using an algorithm, and applied a CGRA schedule to create an efficient application execution on hardware.



**Clockwork** extracted Unified Buffers to enable us to map to the memories on the CGRA.



Thank You



Thank You



Thank You

# Demo notes

## Commands

- `aha map --apps apps/cascade`
- **Generated Log**
  - Full application (CPU enclosed around accelerator = `_hls_target.hw_output`)
  - HalideIR printout of cascade for CGRA
  - Runs CPU and clockwork versions and compares generated output images
  - Runs clockwork memory compilation (with logging in `mem_cout`)
  - Runs MetaMapper

# Demo notes

## Important files

- Halide generator: algorithm+schedule
  - `/aha/Halide-to-Hardware/apps/hardware_benchmarks/apps/cascade/cascade_generator.cpp`
- Generated files from compiler
  - `/aha/Halide-to-Hardware/apps/hardware_benchmarks/apps/cascade/bin/*`
- Input to Clockworks
  - `/aha/Halide-to-Hardware/apps/hardware_benchmarks/apps/cascade/bin/cascade_memory.cpp`
  - `/aha/Halide-to-Hardware/apps/hardware_benchmarks/apps/cascade/bin/cascade_compute.json`
- Output of Clockwork and MetaMapper: mapped design with memories and compute
  - `/aha/Halide-to-Hardware/apps/hardware_benchmarks/apps/cascade/bin/design_top.json`